

## **Assignment #3**

**COMP 424 Big Data**

**Qiangqiang Li (liqiangq@gmail.com)**

## 2 Individual project

**(a). (15 marks) Describe the task including the details of the input data (data source, data size, the original number of features and instances, feature types, whether missing values exist, etc.) and the expected output of the system;**

**Answer:**

- The task is using Decision Tree algorithm to deal with text classification, this is because Decision Trees are a class of very powerful Machine Learning model cable of achieving high accuracy in many tasks while being highly interpretable.
- The input data is Sensorless Drive Diagnosis Data Set, which the detail like the below table, includes features(48) and instances(58509), feature types(Real), whether missing values exist(N/A).

<b>Data source</b>	<a href="https://archive.ics.uci.edu/ml/datasets/dataset+for+sensorless+drive+diagnosis#">https://archive.ics.uci.edu/ml/datasets/dataset+for+sensorless+drive+diagnosis#</a>				
<b>Data Set Characteristics:</b>	Multivariate	<b>Number of Instances:</b>	58509	<b>Area:</b>	Computer
<b>Attribute Characteristics:</b>	Real	<b>Number of Attributes:</b>	49	<b>Date Donated</b>	2015-02-24
<b>Associated Tasks:</b>	Classification	<b>Missing Values?</b>	N/A	<b>Number of Web Hits:</b>	56382

- Features are extracted from electric current drive signals. The drive has intact and defective components. This results in 11 different classes with different conditions. Each condition has been measured several times by 12 different operating conditions, this means by different speeds, load moments and load forces. The current signals are measured with a current probe and an oscilloscope on two phases.
- The expected output of the system is training and testing the dataset, through these 48 features to Predict the different classes which the motor has intact and defective components.

**(b) (5 marks) Describe all the preprocessing steps that are applied to the download data file(s) to obtain the dataset that is used as input of your program, which must be at least 20MB and have at least 30 features.**

**Answer:**

- Sensorless Drive Diagnosis Data Set is from the website, and the original data is about 25M, and using the format of file is text. So need to transform the format to the csv file.
- Because the Features of dataset are number, the good way is using spark mllib .classification lib to transform the JavaRDD, the preprocessing data is format the dataset, the string style translate to double.

- The dataset file (Sensorless\_drive\_diagnosis.csv) include 58509, each data has 48 Features and one label class, and the file size is about 24M .
- **Normalizing data:** use weka tool to preprocessing data, Sensorless\_drive\_diagnosis.csv Transform to Sensorless\_drive\_diagnosis\_normalising.csv. the file size become a smaller one , about 23M.
- **Using PCA:** use weka tool to preprocessing data, Sensorless\_drive\_diagnosis.csv Transform to Sensorless\_drive\_diagnosis\_pca.csv. the file size become a smaller one , about 11M, and feather

**(c) (15 marks) Describe the program using UML class diagrams and/or pseudo-code;**

**Answer: The program using the spark.ml lib and Decision Tree model:**

**1. The initialization of Spark Session, load the dataset and convert TO Java RDD.**

```
SparkSession spark = SparkSession.builder().appName("Decision_Tree").
    //      master("local").
    getOrCreate();

Dataset<Row> kddCSV = spark.read().format("csv").load(data_path);
kddCSV.show();

// convert TO Java RDD
JavaRDD<String> lines = spark.sparkContext().textFile(data_path, 0).toJavaRDD();
JavaRDD<LabeledPoint> linesRDD = lines.map(line -> {
    String[] tokens = line.split(",");
    double[] features = new double[tokens.length - 1];
    for (int i = 0; i < features.length; i++) {
        features[i] = Double.parseDouble(tokens[i]);
    }
    DenseVector v = new DenseVector(features);
    double labelname = Double.parseDouble(tokens[features.length]);
    return new LabeledPoint(labelname, v);
});
```

**2. Index labels, adding meta data to the label column. Fit on whole data set to include all labels in index.**

```
// create the data frame
Dataset<Row> data = spark.createDataFrame(linesRDD, LabeledPoint.class);
data.show();
StringIndexerModel labelIndexer = new StringIndexer().setInputCol("label").setOutputCol("indexedLabel")
    .fit(data);
```

**3. Automatically identify categorical features, and index them, features with > 4 distinct values are treated as continuous.**

```
// Automatically identify categorical features, and index them.
VectorIndexerModel featureIndexer = new
    VectorIndexer().setInputCol("features").setOutputCol("indexedFeatures")
        .setMaxCategories(4) // features with > 4 distinct values are treated as continuous
        .fit(data);
```

#### 4. Split the data into training and test sets (30% held out for testing), Split the data into t Train a Decision Tree model, and the initialization of DecisionTreeClassifier.

```
// Split the data into training and test sets (30% held out for testing)
Dataset<Row>[] splits = data.randomSplit(new double[] { 0.7, 0.3 });
Dataset<Row> training_set = splits[0];
Dataset<Row> test_set = splits[1];

// Train a DecisionTree model
DecisionTreeClassifier dt = new DecisionTreeClassifier().setLabelCol("indexedLabel")
    .setFeaturesCol("indexedFeatures");
```

#### 5. Chain indexers and tree in a Pipeline, and using PipelineModel on training data, Make predictions, Select example rows to display.

```
// Convert indexed labels back to original labels.
IndexToString labelConverter = new
IndexToString().setInputCol("prediction").setOutputCol("predictedLabel")
    .setLabels(labelIndexer.labels());

// Chain indexers and tree in a Pipeline
Pipeline pipeline = new Pipeline()
    .setStages(new PipelineStage[] { labelIndexer, featureIndexer, dt,
labelConverter });

// Train model. This also runs the indexers.
PipelineModel model = pipeline.fit(training_set);

// Make predictions.
Dataset<Row> train_predictions = model.transform(training_set);
Dataset<Row> test_predictions = model.transform(test_set);

// Select example rows to display.
test_predictions.select("predictedLabel", "label", "features").show();

// Select (prediction, true label) and compute test error
MulticlassClassificationEvaluator evaluator = new MulticlassClassificationEvaluator

    .setLabelCol("indexedLabel").setPredictionCol("prediction").setMetricName("accuracy");
```

#### 6. Print the result, include test accuracy, training accuracy, the running time , and the model.

```
double test_accuracy = evaluator.evaluate(test_predictions);
double training_accuracy = evaluator.evaluate(train_predictions);
System.out.println("test error: " + (1.0 - test_accuracy));
System.out.println("test accuracy: " + test_accuracy);
System.out.println("train accuracy: " + training_accuracy);
DecisionTreeClassificationModel treeModel = (DecisionTreeClassificationModel) (model.stages()[2]);
System.out.println("desicion tree model:\n" + treeModel.toDebugString());
long endTime = System.currentTimeMillis();
float running_time = (endTime - startTime) / 1000;
System.out.println("running time: " + running_time + "ms");
```

**(d) (5 marks) Describe how to install and run the program step by step in a Readme.txt file.**

**Answer:** The steps for applying Decision Tree. The contents of Readme.txt is listed below:

**Step 1:** upload java program: Individual\_DecisionTree.java

**Step 2:** ssh to co246a-1.ecs.vuw.ac.nz, and then run environment variables. It was edited a HadoopSetup.csh to avoid inputting the commands manually each time. The contents of HadoopSetup.csh is:

```
◆ export HADOOP_VERSION=2.8.0
◆ export HADOOP_PREFIX=/local/Hadoop/hadoop-$HADOOP_VERSION
◆ export PATH=${PATH}:${HADOOP_PREFIX}/bin
◆ export SPARK_HOME=/local/spark/spark-2.4.2-bin-hadoop2.7
◆ export PATH=${PATH}:${HADOOP_PREFIX}/bin:$SPARK_HOME/bin
◆ export HADOOP_CONF_DIR=$HADOOP_PREFIX/etc/hadoop
◆ export YARN_CONF_DIR=$HADOOP_PREFIX/etc/Hadoop
◆ export
  LD_LIBRARY_PATH=$HADOOP_PREFIX/lib/native:$JAVA_HOME/jre/lib/amd64/s
  erver
```

The spark-hadoop package is already existed in the /local/spark/spark-2.4.2-hadoop2.7 in clusters;

- **Step 3:** upload the dataset: Sensorless\_drive\_diagnosis.csv to one of the Hadoop clusters under /user/<USERNAME>/input. In this case, co246a-1.ecs.vuw.ac.nz is the cluster we used. The command to transfer Sensorless\_drive\_diagnosis.csv is: \$hdfs dfs -put Sensorless\_drive\_diagnosis.csv /user/<USERNAME>/input;
- **Step 4:** copy all the libs from /local/spark/spark-2.4.2-hadoop2.7/jars to the working folder;
- **Step 5:** Go to the working folder, make two new class folder: idt\_class by command: #mkdir idt\_class;
- **Step 6:** patch the java programs to jar patches:  
#javac -cp "libs/\*" -d idt\_class Individual\_DecisionTree.java  
#jar cvf Individual\_DecisionTree.jar idt\_class/ .
- **Step 7:** Submit the work to spark Hadoop jobs:  
#spark-submit --class "nz.ac.vuw.ecs.Individual\_DecisionTree" --master yarn --deploy-mode cluster Individual\_DecisionTree.jar  
/user/<USERNAME>/input/Sensorless\_drive\_diagnosis.csv.

**(e) (15 marks) Compare and discuss the results (including the training and test accuracy, the running time, the model, etc. depending on the program) of the program with and without normalising/scaling data.**

**Answer:**

	Training Acc	Test Acc	Running time	The model	program
<b>Original data</b>	72.84%	73.08%	36ms	DecisionTreeClassificationModel Depth 5 with 39 nodes	nz.ac.vuw.ecs.Individual_DecisionTree.java
<b>normalizing data</b>	81.11%	80.46%	32ms	DecisionTreeClassificationModel Depth 5 with 47 nodes	nz.ac.vuw.ecs.Individual_DecisionTree_N.java

Observably, the scaling data using the same program has a more accurate result than the original data about nearly 9% higher in both training and test dataset. Because using data normalization in order to seek for relation in features, and the normalization data will be better to training and test these features, and give the higher accurate.

On the other hand, the two data sets have the same size, so the running time are the similarly, and use the same model (DecisionTreeClassificationModel), but the result in model has the slight difference, this is may be the seed random.

In summary, the normalization data will be better than the original data to use the Decision Tree algorithm deal with text classification in this case.

**(f) (15 marks) Compare and discuss the results (including the training and test accuracy, the running time, the model, etc. depending on the program) of the program with and without transforming data using PCA.**

	Trainin g Acc	Test Acc	Running time	Data informatio n	The model	program
<b>Original data</b>	72.84%	73.08%	36ms	Size:24M Features:48	Depth 5 with 39 nodes	nz.ac.vuw.ecs.Individual_DecisionTree.java
<b>PCA</b>	52.62%	51.97%	33ms	Size:11M Features:21	Depth 5 with 53 nodes	nz.ac.vuw.ecs.Individual_DecisionTree_PCA.java

Observably, with transforming data using PCA data has a less accurate result than the original data about nearly 20% lower in both training and test dataset. It is mean that using PCA will not suit for this dataset, because PCA is sensitive to the relative scaling of the original variables, the data size become smaller, but the accurate become lower is the mean that these features are relation with each other, the features become less, the means will lack the core mean. It is clear that the transforming data using PCA of the Decision Tree Classification method has not a good way to using in this data.

More over, the running times are the same, but the two data set size have the huge different, and the model (DecisionTreeClassificationModel) using the two ways have the same running time, and the model show the little different in the number of nodes. This is change by the seed random.

In summary, with transforming data using PCA data is worse than the original data to use the Decision Tree algorithm deal with text classification.

## Appendix A. Decision Tree Model for Dataset

[http://co246a-9.ecs.vuw.ac.nz:8088/proxy/application\\_1558496412184\\_1014/](http://co246a-9.ecs.vuw.ac.nz:8088/proxy/application_1558496412184_1014/) **Original**  
[http://co246a-9.ecs.vuw.ac.nz:8088/proxy/application\\_1558496412184\\_1026/](http://co246a-9.ecs.vuw.ac.nz:8088/proxy/application_1558496412184_1026/) **normalising**  
[http://co246a-9.ecs.vuw.ac.nz:8088/proxy/application\\_1558496412184\\_1021/](http://co246a-9.ecs.vuw.ac.nz:8088/proxy/application_1558496412184_1021/) **PCA**

test error: 0.26917839310031777

test accuracy: 0.7308216068996822

train accuracy: 0.7284823284823285

decision tree model:

**DecisionTreeClassificationModel** (uid=dtc\_4805073dfeeb) of depth 5 with 39 nodes

```
If (feature 9 <= -0.051699999999999996)
  If (feature 6 <= -0.0303405)
    If (feature 33 <= 0.0142755)
      Predict: 5.0
    Else (feature 33 > 0.0142755)
      Predict: 4.0
  Else (feature 6 > -0.0303405)
    If (feature 6 <= 0.021951)
      If (feature 6 <= -0.00289445)
        Predict: 4.0
      Else (feature 6 > -0.00289445)
        Predict: 8.0
    Else (feature 6 > 0.021951)
      If (feature 30 <= -0.0141175)
        If (feature 12 <= 0.0022204499999999997)
          Predict: 8.0
        Else (feature 12 > 0.0022204499999999997)
          Predict: 0.0
      Else (feature 30 > -0.0141175)
        Predict: 7.0
  Else (feature 9 > -0.051699999999999996)
    If (feature 10 <= 0.066537)
      If (feature 10 <= -0.0103795)
        If (feature 8 <= -0.0029812)
          If (feature 11 <= -0.021779)
            Predict: 4.0
          Else (feature 11 > -0.021779)
            Predict: 3.0
```



```
Else (feature 8 > -0.0029812)
  If (feature 7 <= 0.0203365)
    Predict: 8.0
  Else (feature 7 > 0.0203365)
    Predict: 7.0
Else (feature 10 > -0.0103795)
  If (feature 10 <= 0.031463500000000005)
    If (feature 8 <= 0.016307500000000003)
      Predict: 1.0
    Else (feature 8 > 0.016307500000000003)
      Predict: 2.0
  Else (feature 10 > 0.031463500000000005)
    Predict: 6.0
Else (feature 10 > 0.066537)
  If (feature 6 <= 0.020324500000000002)
    If (feature 6 <= 0.0025348000000000002)
      If (feature 0 <= -1.995E-5)
        Predict: 4.0
      Else (feature 0 > -1.995E-5)
        Predict: 3.0
    Else (feature 6 > 0.0025348000000000002)
      Predict: 6.0
  Else (feature 6 > 0.020324500000000002)
    If (feature 27 <= 0.022981)
      Predict: 9.0
    Else (feature 27 > 0.022981)
      If (feature 4 <= -1.145E-5)
        Predict: 6.0
      Else (feature 4 > -1.145E-5)
        Predict: 9.0
```

running time: 36.0ms

---

test error: 0.19532449803765428

test accuracy: 0.8046755019623457

train accuracy: 0.8111317435496481

desision tree model:

DecisionTreeClassificationModel (uid=dtc\_bd954f952a59) of depth 5 with 47 nodes

```
If (feature 11 <= 0.4954545)
If (feature 10 <= 0.364080500000000003)
If (feature 7 <= 0.650576)
If (feature 9 <= 0.2970755)
If (feature 6 <= 0.5088895)
    Predict: 5.0
Else (feature 6 > 0.5088895)
    Predict: 4.0
Else (feature 9 > 0.2970755)
If (feature 11 <= 0.344108500000000004)
    Predict: 4.0
Else (feature 11 > 0.344108500000000004)
    Predict: 3.0
Else (feature 7 > 0.650576)
If (feature 6 <= 0.765679500000000001)
If (feature 24 <= 0.3585205)
    Predict: 0.0
Else (feature 24 > 0.3585205)
    Predict: 8.0
Else (feature 6 > 0.765679500000000001)
If (feature 10 <= 0.3408775)
    Predict: 7.0
Else (feature 10 > 0.3408775)
    Predict: 0.0
Else (feature 10 > 0.364080500000000003)
If (feature 11 <= 0.4348355)
If (feature 8 <= 0.73367549999999999)
If (feature 6 <= 0.5769925)
    Predict: 3.0
Else (feature 6 > 0.5769925)
    Predict: 1.0
Else (feature 8 > 0.73367549999999999)
If (feature 8 <= 0.8016485)
    Predict: 10.0
Else (feature 8 > 0.8016485)
    Predict: 2.0
Else (feature 11 > 0.4348355)
If (feature 11 <= 0.454915500000000003)
If (feature 33 <= 0.646206000000000001)
```

```
Predict: 6.0
Else (feature 33 > 0.6462060000000001)
  Predict: 10.0
Else (feature 11 > 0.45491550000000003)
  Predict: 6.0
Else (feature 11 > 0.4954545)
  If (feature 6 <= 0.7724614999999999)
    If (feature 18 <= 0.33653500000000003)
      Predict: 9.0
    Else (feature 18 > 0.33653500000000003)
      If (feature 6 <= 0.657035)
        If (feature 0 <= 0.702442)
          Predict: 4.0
        Else (feature 0 > 0.702442)
          Predict: 3.0
      Else (feature 6 > 0.657035)
        Predict: 6.0
    Else (feature 6 > 0.7724614999999999)
      If (feature 27 <= 0.3949355)
        If (feature 6 <= 0.7879290000000001)
          If (feature 30 <= 0.3612145)
            Predict: 9.0
          Else (feature 30 > 0.3612145)
            Predict: 6.0
        Else (feature 6 > 0.7879290000000001)
          Predict: 9.0
      Else (feature 27 > 0.3949355)
        If (feature 6 <= 0.7945795)
          Predict: 6.0
        Else (feature 6 > 0.7945795)
          Predict: 9.0
```

running time: 32.0ms

=====

test error: 0.46218010554389155

test accuracy: 0.5378198944561084

train accuracy: 0.5375434133933376

desision tree model:

DecisionTreeClassificationModel (uid=dtc\_87e3cfb7372e) of depth 5 with 51 nodes

```
If (feature 1 <= 2.17389450000000003)
  If (feature 1 <= -2.9737135)
    If (feature 2 <= 1.4098325)
      If (feature 18 <= 0.75677250000000001)
        Predict: 9.0
      Else (feature 18 > 0.75677250000000001)
        If (feature 2 <= -0.178970000000000002)
          Predict: 9.0
        Else (feature 2 > -0.178970000000000002)
          Predict: 10.0
    Else (feature 2 > 1.4098325)
      If (feature 3 <= -1.8251435)
        If (feature 19 <= -0.12217449999999999)
          Predict: 6.0
        Else (feature 19 > -0.12217449999999999)
          Predict: 10.0
      Else (feature 3 > -1.8251435)
        If (feature 18 <= 0.4248865)
          Predict: 9.0
        Else (feature 18 > 0.4248865)
          Predict: 10.0
    Else (feature 1 > -2.9737135)
      If (feature 1 <= -0.389029)
        If (feature 19 <= -0.3500835)
          If (feature 18 <= 0.371115)
            Predict: 6.0
          Else (feature 18 > 0.371115)
            Predict: 10.0
        Else (feature 19 > -0.3500835)
          If (feature 19 <= 0.6628535)
            Predict: 2.0
          Else (feature 19 > 0.6628535)
            Predict: 7.0
      Else (feature 1 > -0.389029)
        If (feature 19 <= -0.06562599999999999)
          If (feature 1 <= 0.32309)
            Predict: 1.0
          Else (feature 1 > 0.32309)
```

```
Predict: 3.0
Else (feature 19 > -0.06562599999999999)
  If (feature 19 <= 0.58487050000000001)
    Predict: 8.0
  Else (feature 19 > 0.58487050000000001)
    Predict: 7.0
Else (feature 1 > 2.17389450000000003)
  If (feature 1 <= 5.73697450000000005)
    If (feature 18 <= -0.444182500000000004)
      If (feature 11 <= -0.67647499999999999)
        If (feature 0 <= -1.21727)
          Predict: 3.0
        Else (feature 0 > -1.21727)
          Predict: 5.0
      Else (feature 11 > -0.67647499999999999)
        If (feature 19 <= 0.157006)
          Predict: 4.0
        Else (feature 19 > 0.157006)
          Predict: 5.0
    Else (feature 18 > -0.444182500000000004)
      If (feature 11 <= 1.237725)
        If (feature 0 <= -5.205881)
          Predict: 4.0
        Else (feature 0 > -5.205881)
          Predict: 5.0
      Else (feature 11 > 1.237725)
        If (feature 19 <= -0.4784385)
          Predict: 4.0
        Else (feature 19 > -0.4784385)
          Predict: 5.0
    Else (feature 1 > 5.73697450000000005)
      If (feature 0 <= 6.8679745)
        Predict: 5.0
      Else (feature 0 > 6.8679745)
        If (feature 1 <= 6.71557950000000005)
          Predict: 3.0
        Else (feature 1 > 6.71557950000000005)
          Predict: 9.0
```

running time: 33.0ms